

Recurrent Neural Networks (RNNs) - A Brief Overview

What is an RNN?

A Recurrent Neural Network (RNN) is a type of neural network designed for processing sequential data, such as time series data.

Key Concepts:

- Hidden State: Each RNN cell maintains a hidden state vector $h(t)$ that captures information from previous steps.
- Sequence Processing: Inputs are processed one at a time (e.g., word by word in a sentence), and the network's output depends on the current input and the hidden state.
- Parameter Sharing: RNNs use the same weights at each time step, reducing the number of parameters.

Mathematics Behind RNNs:

For input $x(t)$ and hidden state $h(t)$ at time step t :

$$h(t) = \tanh(W_{xh} * x(t) + W_{hh} * h(t-1) + b_h)$$

$$y(t) = W_{hy} * h(t) + b_y$$

Where:

- W_{xh} , W_{hh} , W_{hy} are weight matrices
- b_h , b_y are bias vectors
- \tanh is an activation function

Problems with Vanilla RNNs:

- Vanishing and Exploding Gradients: Hard to capture long-term dependencies.
- Short Memory: Struggles to remember information from far back in the sequence.

Variants of RNNs:

- LSTM (Long Short-Term Memory): Uses gates to control information flow.
- GRU (Gated Recurrent Unit): A simplified, efficient version of LSTM.

Applications:

- Natural Language Processing (translation, sentiment analysis)
- Time Series Forecasting
- Speech Recognition
- Music Generation